

Basics of Operating Systems (2)

Dr. Qin Zhou

Coventry University
email: q.zhou@coventry.ac.uk

Files and directories

- Files are stored on disks, CDs etc.
 - these devices store data in fixed-size blocks
- Different devices use different block sizes
- Different devices have different capabilities
 - layout of data blocks may effect efficiency
- Operating system need to hide these device-dependent aspects

Files and directories

- The filesystem lets you refer to files by name
 - the filesystem defines a *namespace* (a set of names, in the same way that an address space is a set of addresses)
- Where the files are located becomes irrelevant
 - the filesystem can organise data layout to optimise access to the data

Files and directories

- Filesystems store information in *volumes*
 - usually a single device, e.g. a disk or CD
 - a single devices can be *partitioned* into multiple volumes
 - a single volume might span multiple devices
- Each volume is treated as a numbered sequence of blocks
 - filesystem needs to keep track of file names and corresponding block numbers


Files and directories

- Each volume is treated as a numbered sequence of blocks
 - filesystem needs to keep track of file names and corresponding block numbers
- A list of filenames and locations is called a *directory*
 - this is data that can be stored as just another file
 - there must be a *root directory* at a known place on each volume

Files and directories

- By treating directories as (special) files, there can be multiple directories on a volume
 - directories can contain entries for files and other (sub-) directories
 - this leads to a tree-structured *hierarchical filesystem*

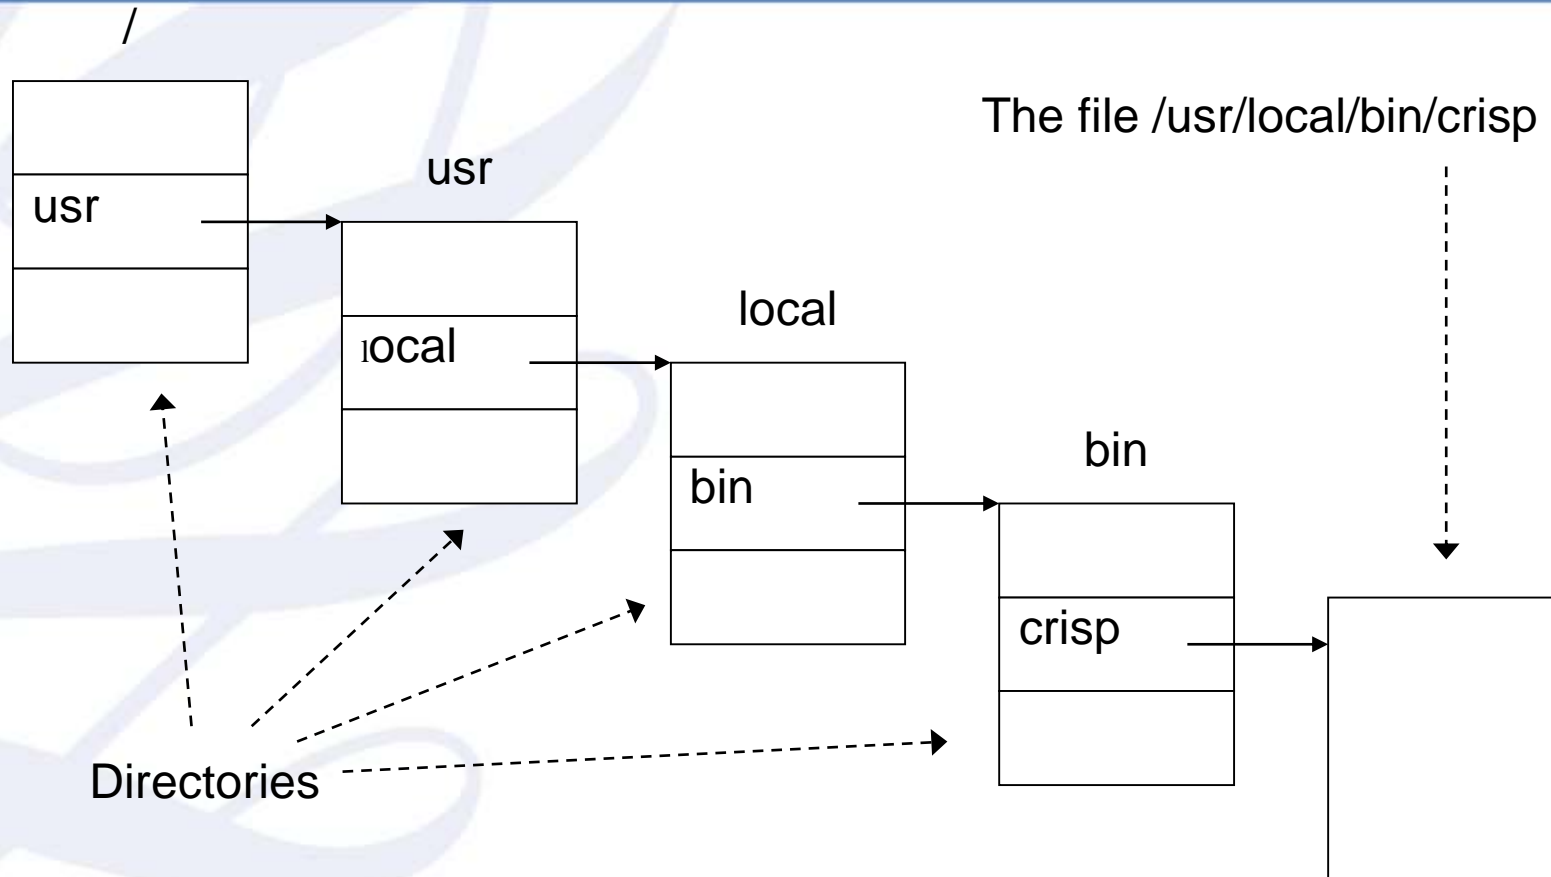
Files and directories



Pathnames

- To identify a file uniquely, we need to use a *pathname*
 - this specifies the path from the root directory through any subdirectories to the file
- Unix naming conventions:
 - the root directory is called ‘/’
 - ‘/’ is also used to separate individual directory and file names in a pathname

Pathnames



Pathnames

- Most filesystems keep track of a *current directory*
 - names which do not begin with '/' (the root directory) are taken to be relative to the current directory
- Special names:
 - '.' is the current directory
 - '..' is the *parent directory* which contains the current directory

Pathnames

- If the current directory is `/usr/local/bin`:
 - the file `/usr/local/bin/crisp` can be referred to as `crisp` or as `./crisp`
 - the directory `/usr/local/bin` can be referred to as `.`
 - the directory `/usr/local` can be referred to as `..`
 - the file `/usr/local/man/crisp.1` can be referred to as `../man/crisp.1`

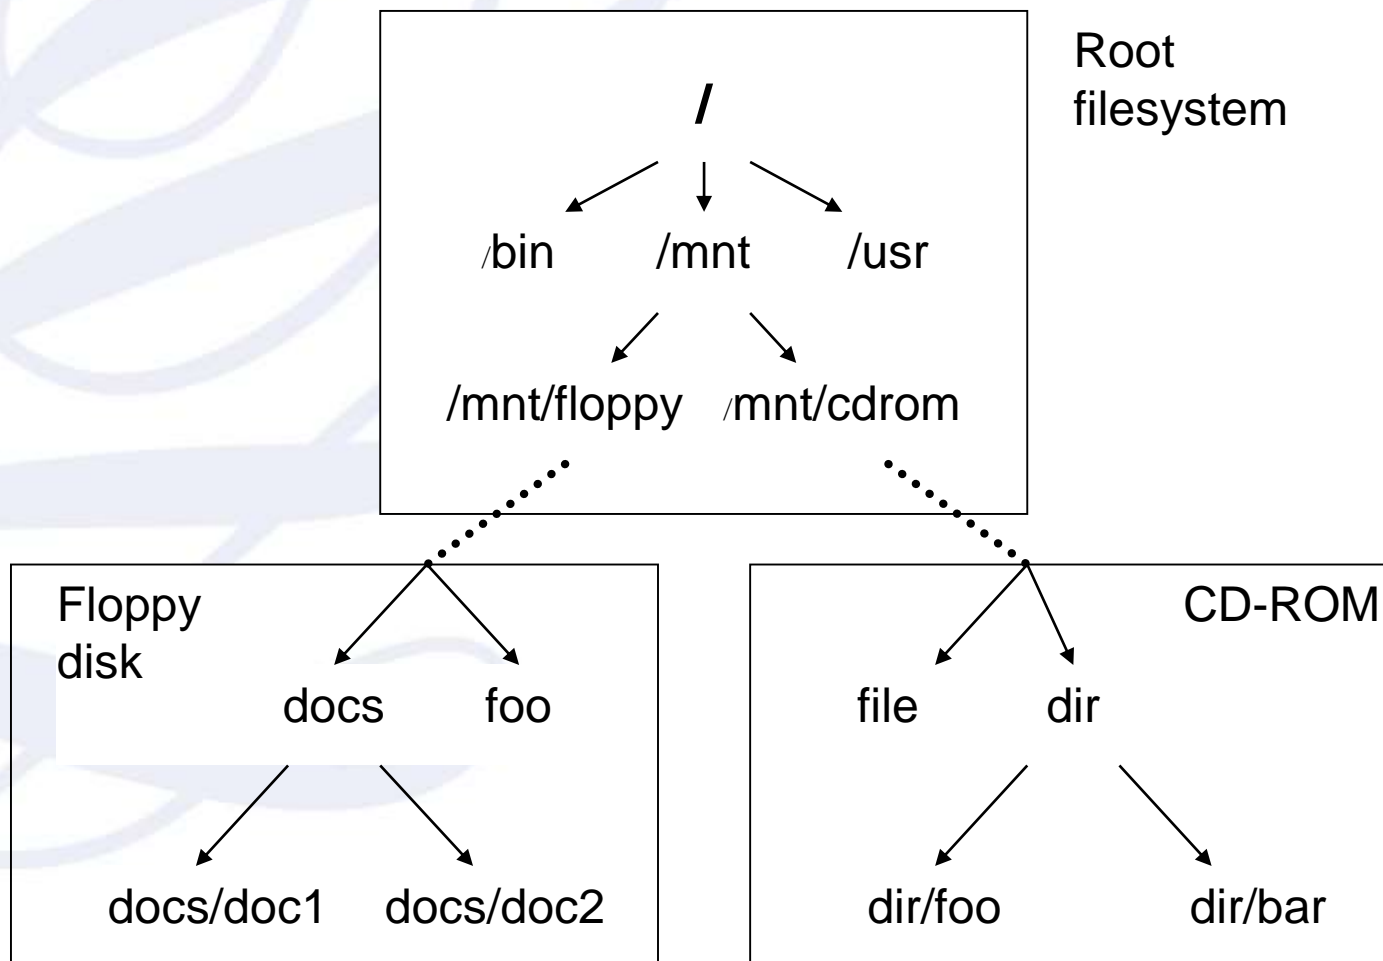
Windows pathnames

- Same as Unix, but:
 - ‘\’ used instead of ‘/’
 - names must be preceded by a volume identifier, e.g. ‘C:’ for the primary hard disk
- ‘Drive letters’ like C: limits system to 26 volumes
 - disks from remote machines can be *mapped* to drive letters
 - UNC names: refer to files on remote machine using \\machine as a prefix instead of e.g. C:

Multiple volumes in Unix

- Unix uses a unified filesystem
- There is a *root filesystem* (initially a single volume)
 - additional volumes are *mounted* as subdirectories within the root filesystem
 - the mounted volume's directory structure replaces the directory it is mounted onto

Multiple volumes in Unix



Multiple volumes in Unix

- Different volumes will be organised differently
 - the *virtual filesystem* (VFS) switches between different filesystems as *mount points* are encountered when tracing through a pathname
- Volumes must be explicitly unmounted before removal

The virtual filesystem


`open ("/mnt/cdrom/dir/foo", ...)`

Virtual file system:

Real file systems:

Device drivers:

Devices:



Types of files

- VMS (ancestor of Windows NT) provided different file types:
 - sequential: could process from start to end
 - direct access: fixed-size records, could process in any sequence by record number
 - indexed: similar, but records identified by a symbolic key
- This makes it hard to write a program to process any file

Types of file

- Still need to distinguish different types of file according to content
 - e.g. text files vs. executable binary files
- Both Unix and Windows use suffixes (extensions) to identify file contents:
 - e.g. file.txt or file.exe
 - Windows associates different actions with different extensions
 - Unix extensions are just a naming convention

Types of file

- Executable programs are identified:
 - by their extension in Windows (.com, .exe, .bat or .cmd)
 - by a file attribute in Unix
- Unix shell scripts can be made executable
 - first line ('shebang' line) specifies pathname of program to use to process file

```
!#/bin/sh
```


Filenames

- Limited character set
 - characters like ‘*’ are interpreted specially and should not be used
 - names including spaces must be quoted (“...”)
- Limited length
 - MS-DOS: 8 characters + 3 for extension
 - early Unix systems: 14 characters
 - now: 255 characters

Links and shortcuts

- It is sometimes useful to keep the same file in several places
 - keep copies? but then they need to be re-copied after any changes...
- Unix uses *inodes* to hold all information about file layout on disk etc.
 - directory entries refer to inodes
 - multiple directories can contain *links* to the same inode

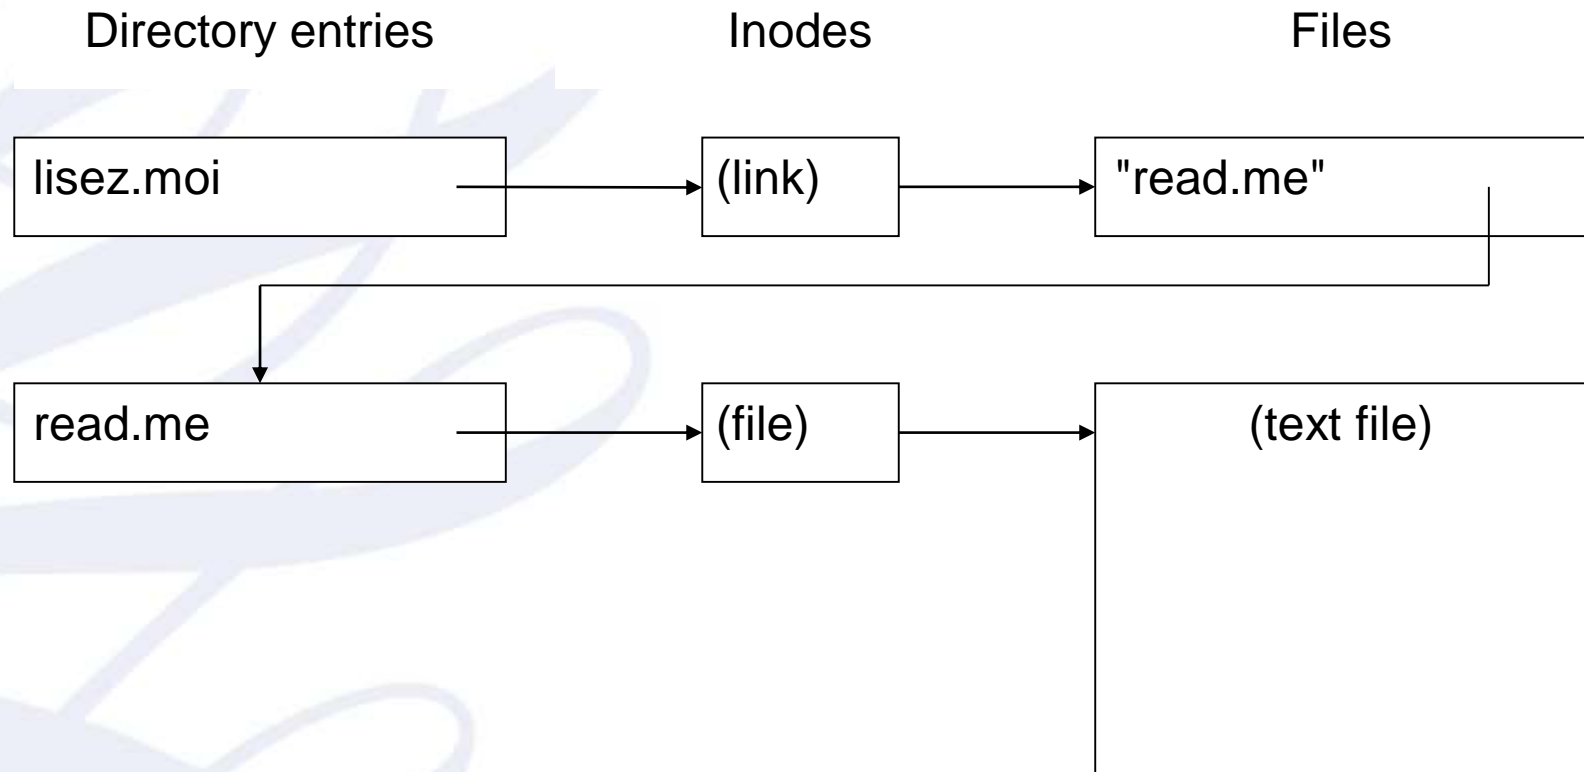
Links and shortcuts



Links and shortcuts

- Links can only refer to files on the same volume
 - solution: *soft links* which are files containing the pathname of the file being linked to
 - file attribute used to mark file as a soft link
- Windows provides *shortcuts*
 - contains more information than just the name
 - not exactly the same: shell accesses real file when icon activated, but otherwise a shortcut is a normal file

Links and shortcuts



File attributes

- Apart from name, directory entry (or inode) should provide some or all of:
 - timestamps (creation, last update, last access)
 - ownership (individual and group)
 - file size (bytes)
 - file type (e.g. file or directory)
 - access permissions
- Other attributes possible (version number, password, ...)

Access permissions

- Permissions include read, write, execute, delete, append, ...
- Unix: each user belongs to one or more groups
 - separate permissions for file owner, owning group and others (read, write, execute)
- Windows: each file has an access control list
 - separate permissions for individuals or groups

Example file systems

- FAT (MS-DOS, Windows)
 - other systems can also use this as a ‘lowest common denominator’
 - used for floppy disks
- Ext 2/3/4 (Linux filesystem’)
 - based on original BSD Unix filesystem
- NTFS (Windows NT, 2000, XP, Vista, 7/8)

Reference

